References

Week 1:

- The Sequential model (TensorFlow Documentation)
- The Functional API (TensorFlow Documentation)

Week 2:

- Deep Residual Learning for Image Recognition (He, Zhang, Ren & Sun, 2015)
- deep-learning-models/resnet50.py/ (GitHub: fchollet)
- MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications (Howard, Zhu, Chen, Kalenichenko, Wang, Weyand, Andreetto, & Adam, 2017)
- MobileNetV2: Inverted Residuals and Linear Bottlenecks (Sandler, Howard, Zhu, Zhmoginov &Chen, 2018)
- EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks (Tan & Le, 2019)

Week 3:

- You Only Look Once: Unified, Real-Time Object Detection (Redmon, Divvala, Girshick & Farhadi, 2015)
- YOLO9000: Better, Faster, Stronger (Redmon & Farhadi, 2016)
- YAD2K (GitHub: allanzelener)
- YOLO: Real-Time Object Detection
- Fully Convolutional Architectures for Multi-Class Segmentation in Chest Radiographs (Novikov, Lenis, Major, Hladůvka, Wimmer & Bühler, 2017)
- Automatic Brain Tumor Detection and Segmentation Using U-Net Based Fully Convolutional Networks (Dong, Yang, Liu, Mo & Guo, 2017)
- U-Net: Convolutional Networks for Biomedical Image Segmentation (Ronneberger, Fischer & Brox, 2015)

Week 4:

- FaceNet: A Unified Embedding for Face Recognition and Clustering (Schroff, Kalenichenko & Philbin, 2015)
- DeepFace: Closing the Gap to Human-Level Performance in Face Verification (Taigman, Yang, Ranzato & Wolf)
- facenet (GitHub: davidsandberg)
- How to Develop a Face Recognition System Using FaceNet in Keras (Jason Brownlee, 2019)
- keras-facenet/notebook/tf_to_keras.ipynb (GitHub: nyoki-mtl)
- A Neural Algorithm of Artistic Style (Gatys, Ecker & Bethge, 2015)
- Convolutional neural networks for artistic style transfer
- TensorFlow Implementation of "A Neural Algorithm of Artistic Style"

- Very Deep Convolutional Networks For Large-Scale Image Recognition (Simonyan & Zisserman, 2015)
- Pretrained models (MatConvNet)